R³: Resonance-Based Relational Reasoning

A Unified White-Box Architecture for Interpretable Audio Representation
and Neuroscience-Grounded Cognitive Synthesis

Amaç Erdem
Boston University, College of Fine Arts
amac.erdem.ae@gmail.com

Abstract—We present **R³** (Resonance-Based Relational Reasoning), a unified white-box architecture for interpretable audio representation and neuroscience-grounded synthesis. R³ provides *full interpretability* through a 256D representation: 128D core ontology (L_0-L_9) plus 128D cross-layer relations in three temporal tiers, extending to 6,144D via τ-Morph.
The framework integrates with $\mathbf{C^0}$ (Cognitive Meta-Supervisor), grounded in 387 neuroimaging studies (1,116 claims, $d = 2.117$). Key innovations: **Interpretability Gradient**, **OSE** with zero-entanglement bijection, and **LAM** for neuroscience-grounded supervision.
R³-Core (48,000 lines, 193 modules) is in **active H200 GPU training**, achieving 149× real-time (RTF=0.0067) with 91% benchmark validation.

Keywords: interpretable AI, audio representation, neural codec, psychoacoustics, cognitive neuroscience

1. Introduction

The past decade has witnessed remarkable advances in neural audio processing. Models like Jukebox [1], MusicLM [2], and MusicGen [3] generate convincing musical audio. Neural codecs such as EnCodec [4] and RAVE [5] achieve impressive compression. However, these systems share a fundamental limitation: their latent representations are *black boxes*.

The Translation Problem

Consider a jazz drummer creating a "swing feel." Four experts offer four vocabularies:

- **Music theorist**: "Syncopation with 2:1 subdivision"
- **Audio engineer**: "10-50 ms microtiming deviations"
- **Neuroscientist**: "2 Hz delta band entrainment"
- **Performer**: "It swings"

These descriptions refer to the *same phenomenon* using different coordinate systems. This **translation problem** manifests in: (1) no interpretability, (2) no controllability, (3) no verifiability, and (4) no transferability.

The White-Box Alternative

We propose **R³**, a white-box alternative where every latent dimension has explicit semantic meaning. R³ is grounded in over a century of psychoacoustic research—from Helmholtz [6] to Bregman [7] to Tymoczko [8]. Rather than learning representations from scratch, R³ *extracts* dimensions known to be perceptually meaningful.

Production Status

This is not a theoretical proposal. R³-Core is a complete implementation currently in active training:

- **Codebase**: 48,000 lines, 193 PyTorch modules
- **Training**: Active on NVIDIA H200 GPUs (cloud)
- **Performance**: 149× real-time (RTF=0.0067)
- **Validation**: 91% benchmark pass rate

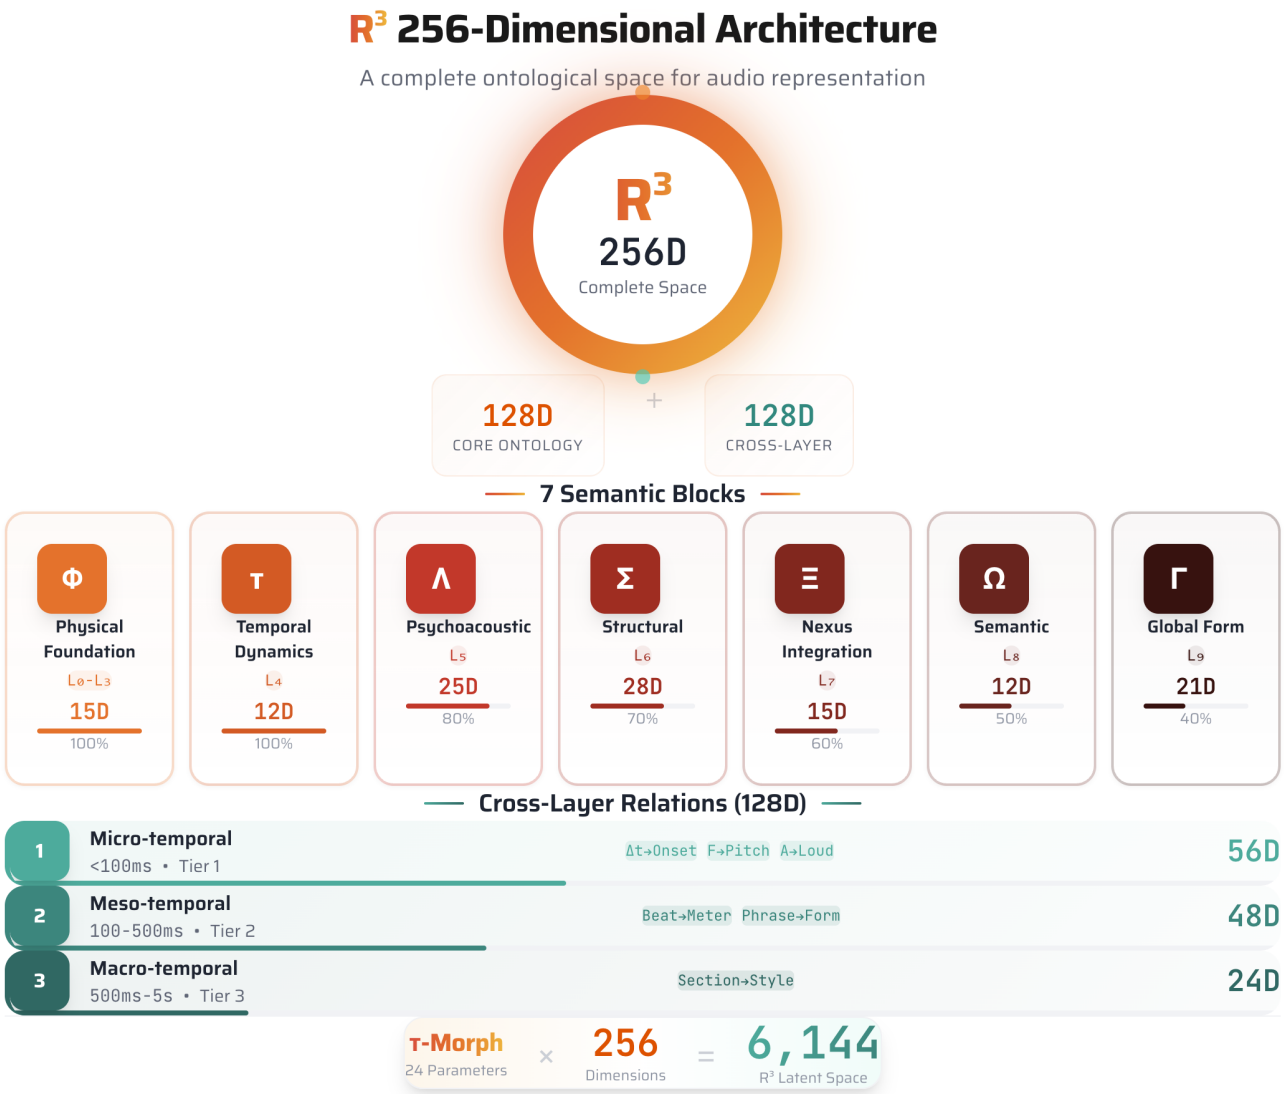

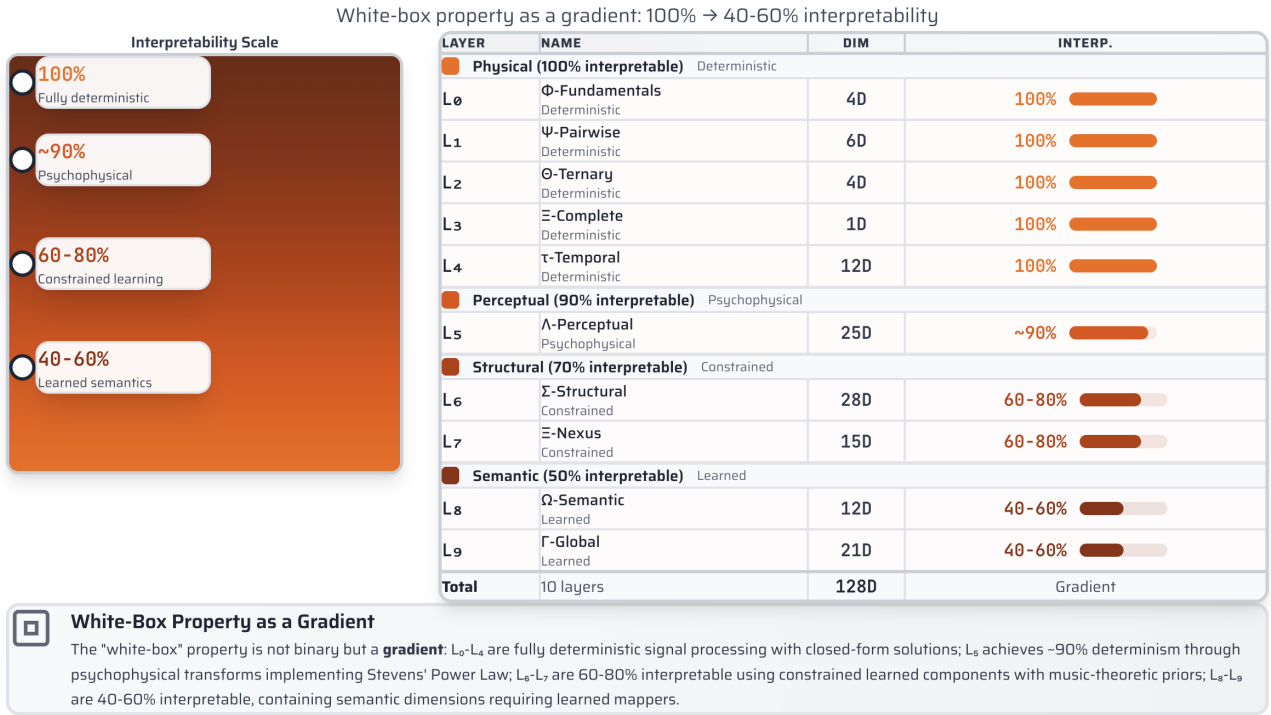

2. Related Work

Neural Audio Codecs

Neural audio compression has advanced rapidly since WaveNet [9] demonstrated high-quality autoregressive synthesis. VQ-VAE [10] introduced discrete latent codes for audio. EnCodec [4] achieves near-transparent quality at 6 kbps using residual vector quantization. DAC [11] improves compression with enhanced discriminators. RAVE [5] enables real-time synthesis. All these systems learn latent representations end-to-end; dimensions lack semantic meaning.

Music Understanding Models

Large-scale music understanding has progressed through contrastive learning. MuLan [12] learns joint audio-text embeddings. MusicLM [2] extends this for generation. MusicGen [3] provides efficient single-stage synthesis. AudioLDM [13] applies latent diffusion. These models operate in 512-1024D spaces, but dimensions lack individual meaning.

Interpretable Audio Representations

Classical features (MFCCs, chroma, spectral descriptors) are interpretable but limited. Recent work on disentangled representations [14] attempts separation, but interpretation remains challenging. Our approach differs: rather than hoping disentanglement emerges, we *define* dimensions based on established psychoacoustic knowledge.

Psychoacoustics and Music Cognition

R³ draws on Helmholtz's theory of beats [6], Tenney's historical analysis [15], Pressnitzer and McAdams' spectral music research [16], Bregman's auditory scene analysis [7], and Tymoczko's geometric foundations [8].

3. R³ Architecture

R³ is built on three principles: (1) *interpretability first*—every dimension corresponds to a defined property; (2) *psychoacoustic grounding*—dimensions reflect human perception [6], [7], [8]; and (3) *hierarchical organization*—lower layers capture physical properties, higher layers capture cognitive properties.

256D Representation Structure

The complete R³ representation is:

$$\mathbf{r} \in \mathbb{R}^{256} = \underbrace{\mathbf{r}_{\text{core}}}_{\text{128D: L}_0\text{-L}_9} \oplus \underbrace{\mathbf{r}_{\text{cross}}}_{\text{128D: 16 relations}} \tag{1}$$

L_0-L_9 Core Ontology (128D)

The ontology spans 10 layers with an **Interpretability Gradient**:

The "white-box" property is a *gradient*: L_0-L_4 are fully deterministic signal processing; L_5-L_7 use constrained learned components; L_8-L_9 require learned semantic mappers.

Layer	Name	Dim	Interp.
L_0	Φ-Fundamentals	4	100%
L_1	Ψ-Pairwise	6	100%
L_2	Θ-Ternary	4	100%
L_3	Ξ-Complete	1	100%
L_4	τ-Temporal	12	100%
L_5	Λ-Perceptual	25	∼90%
L_6	Σ-Structural	28	60-80%
L_7	Ξ-Nexus	15	60-80%
L_8	Ω-Semantic	12	40-60%
L_9	Γ-Global	21	40-60%

Table 1: L_0-L_9 core ontology with interpretability gradient.

Cross-Layer Relations (128D)

Beyond individual layers, R³ captures 16 psychoacoustically meaningful layer interactions organized into three temporal tiers:

Tier 1 (<100ms, 56D): Micro-temporal—cochlear mapping, Stevens' power law [17], critical band roughness, spectral brightness.

Tier 2 (100-500ms, 48D): Meso-temporal—frequency streaming, attention capture, the critical $L_4 \leftrightarrow L_5$ *Flux→Rhythm* relation.

Key Innovation: Delta-Band Filter. Neural evidence [18] shows cortical synchronization is strongest for *spectral flux*, not amplitude envelope. The X10 cross-layer explicitly filters for ∼2 Hz (≈120 BPM), aligning acoustic processing with basal ganglia-thalamo-cortical beat perception.

Tier 3 (500ms-5s, 24D): Macro-temporal—melodic contour, harmonic consonance, tonal center estimation.

τ-Morph: 24D Temporal Expansion

Beyond instantaneous 256D features, musical meaning resides in *how* quantities evolve. Given node trajectory $U \in \mathbb{R}^{W \times K}$ across W windows, we define six 4D groups (total 24D):

Central (4D): Global mean $\tau_1 = \mu$, median τ_2, time centroid $\tau_3 = \sum_{w,k} w \cdot u_{w,k} / \sum u_{w,k}$, feature centroid τ_4.

Shape (4D): Skewness $\tau_5 = \frac{\mathbb{E}[(v-\bar{v})^3]}{\mathbb{E}[(v-\bar{v})^2]^{3/2}}$, kurtosis τ_6, normalized entropy $\tau_7 = -\sum_b p_b \log p_b / \log B$, spectral flatness τ_8.

Dynamics (4D): Range $\tau_9 = \max u - \min u$, mean variance τ_{10}, IQR τ_{11}, mean gradient $\tau_{12} = \frac{1}{(W-1)K} \sum_{w,k} |\Delta u_{w,k}|$.

Frequency (4D): From temporal PSD P_ω: dominant modulation freq $\tau_{13} = \arg\max_\omega P_\omega$, bandwidth τ_{14}, 90% rolloff τ_{15}, modulation flux τ_{16}.

Complexity (4D): Lempel-Ziv τ_{17}, permutation entropy τ_{18}, Hurst exponent τ_{19}, fractal dimension τ_{20}.

Position (4D): Onset τ_{21}, offset τ_{22}, peak location τ_{23}, center of mass τ_{24}.

This yields the **R³ Latent Space**: $\mathbf{z}_w = \tau\text{-Morph}(\mathbf{r}_{w_1}, \ldots, \mathbf{r}_{w_T}) \in \mathbb{R}^{256 \times 24} = \mathbb{R}^{6,144}$ per window. For a 3-minute piece with 100 windows and

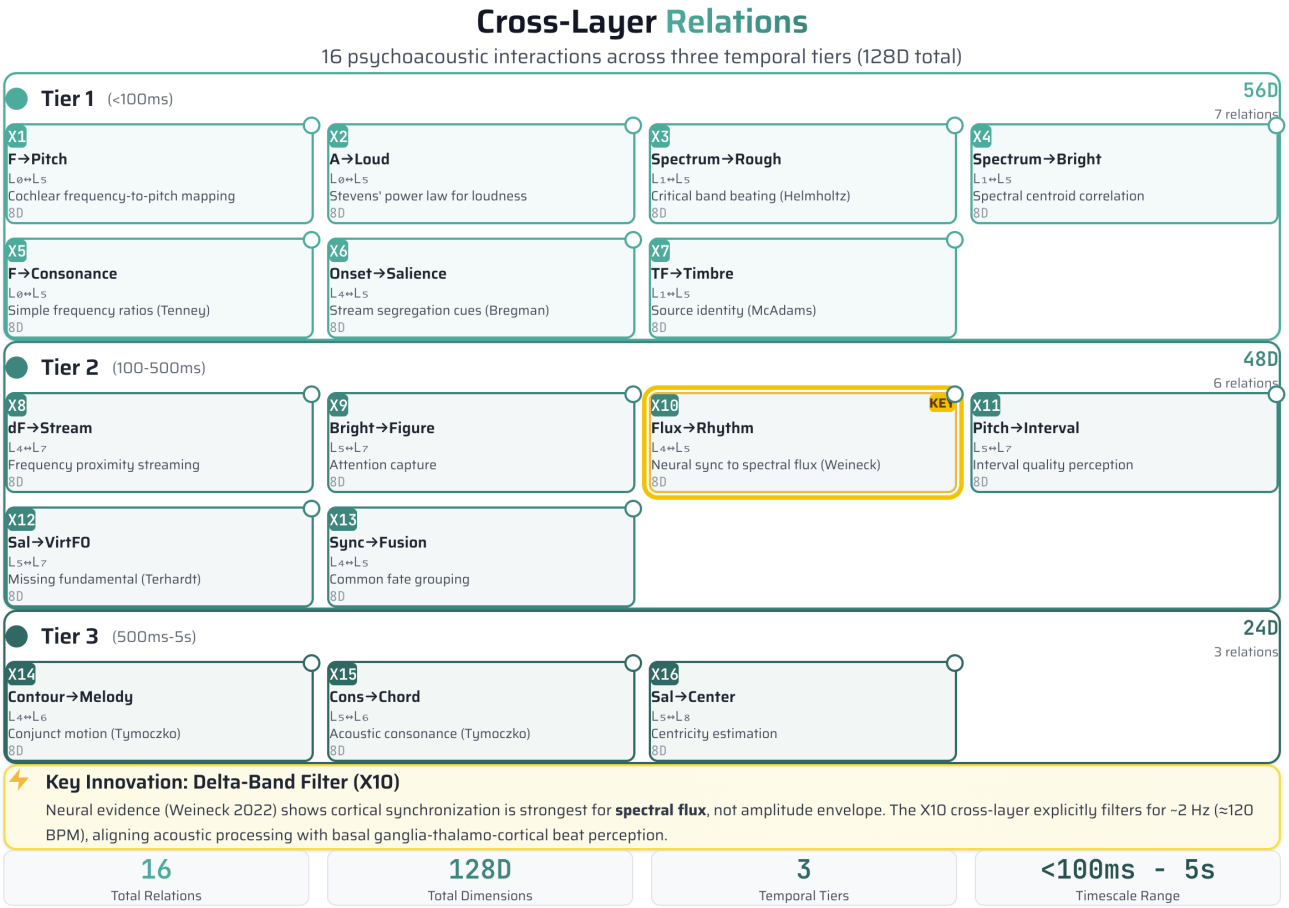

5 sections: Piece Total $\approx 616,656$D—dwarfing neural codecs (MusicLM: 1024D) while maintaining full interpretability.

Ξ-Nexus: Relational Reasoning

The Ξ-Nexus module performs graph-based relational reasoning over τ-Morph nodes via sparse graph $\mathscr{G} = (\mathscr{V}, \mathscr{E})$ with five edge types:

1. **Temporal** ($\mathscr{E}_T$): Adjacent time windows
2. **Frequency** ($\mathscr{E}_F$): Neighboring frequency bands
3. **Cross-scale** ($\mathscr{E}_C$): Higher layers to aggregates
4. **Harmonic** ($\mathscr{E}_H$): Frequency ratios $\in \{2,3,4,5,6\}$
5. **Phase** ($\mathscr{E}_P$): Phase-locking value $> \theta_P$

Sparse Relational Attention. For each relation $r \in \{T, F, C, H, P\}$:

$$\alpha_{ij}^{(r)} = \text{softmax}\left(\frac{(W_r^Q h_i)^\top (W_r^K h_j)}{\sqrt{d_h}}\right) \tag{2}$$

$$h_i^{(l+1)} = h_i^{(l)} + \sigma\left(\sum_r \gamma_r \sum_{j \in \mathscr{N}_r(i)} \alpha_{ij}^{(r)} W_r^V h_j\right) \tag{3}$$

where γ_r are learnable relation weights. Harmonic edges capture chord progressions; phase edges capture rhythmic synchronization.

Hierarchical Pooling:

$$g_w = \text{Pool}_n(h_{w,n}^{(L)}) \quad \text{(window-level)} \tag{4}$$

$$s_\ell = \text{Pool}_{w \in \mathscr{S}_\ell}(g_w) \quad \text{(section-level)} \tag{5}$$

$$u = \text{Pool}_\ell(s_\ell) \quad \text{(piece-level)} \tag{6}$$

This hierarchy provides g_w for local MIR tasks (beat detection), s_ℓ for form analysis, and u for global classification and T_3 conditioning.

4. Dual-Use Framework

R^3 functions both as an analytical tool (Audio $\to R^3$) and a generative latent space ($R^3 \to$ Audio).

R^3 Encoder

The encoder extracts R^3 features via hierarchical signal processing: STFT decomposition [19], fundamental extraction (L_0-L_3), temporal derivatives (L_4), perceptual transforms (L_5) using Stevens' power law, structural analysis (L_6-L_7), semantic inference (L_8-L_9), and cross-layer relation computation.

OSE Decoder: Ontological Synthesis Engine

The **Ontological Synthesis Engine (OSE)** implements 100% interpretable synthesis: $\mathscr{D}(\mathbf{r}) = \mathscr{S}_0 \circ \mathscr{S}_1 \circ \cdots \circ \mathscr{S}_9(\mathbf{r})$ where each $\mathscr{S}_i$ is a deterministic layer-wise operator.

Layer-Wise Synthesis ($L_9 \to L_0$):

- **High-Level Planning ($L_9 \to L_6$)**: Form envelope $E_{\text{form}}(t)$ from Γ_{section}, harmonic skeleton $\mathbf{C}(t)$ via voice-leading
- **Perceptual Targets ($L_5 \to L_4$)**: Stevens' Power Law inversion: $f_0(t) = \text{inv_mel}(\Lambda_{\text{pitch}}(t))$, $A(t) = \Lambda_{\text{loudness}}(t)^{1/0.30}$
- **Spectral Construction ($L_3 \to L_1$)**: $X(t,f) = \Theta_{TFA}(t,f) \cdot e^{j \cdot \Theta_{TF\phi}(t,f)}$
- **Waveform Synthesis (L_0)**: Inverse STFT or additive: $x(n) = \sum_{k=1}^{K} A_k(n) \cdot \sin(2\pi f_k n + \phi_k(n))$

Zero-Entanglement Bijection. For every synthesis parameter p, there exists a unique R^3 dimension r_d:

$$\frac{\partial p}{\partial r_d} \neq 0 \quad \text{and} \quad \frac{\partial p}{\partial r_{d'}} = 0 \quad \forall d' \neq d \tag{7}$$

This guarantees modifying one dimension affects exactly one synthesis parameter—impossible in black-box decoders.

Tiered Quality Model:

- **Tier A (Harmonic)**: Pitched instruments, vocals—OSE native, high fidelity (>20 dB SI-SDR)
- **Tier B (Mixed)**: Transient-rich sounds—OSE + residual noise model
- **Tier C (Complex)**: Percussion, noise-based—neural refinement layer recommended

Semantic Consistency: $\forall x : d(\text{Synthesize}(\text{Analyze}(x)), x) \leq \varepsilon$

5. C^0: Cognitive Supervision Framework

C^0 (Cognitive Meta-Supervisor) provides neuroscience-grounded supervision derived from comprehensive meta-analysis of **387 neuroimaging studies** (448 extractions, 1,116 empirical claims, $N \approx 7,235$ participants).

Meta-Analysis Methodology

We employed DerSimonian-Laird random-effects meta-analysis following PRISMA 2020 guidelines:

$$\hat{\delta} = \frac{\sum_{i=1}^{k} w_i d_i}{\sum_{i=1}^{k} w_i}, \quad w_i = \frac{1}{\sigma_i^2 + \hat{\tau}^2} \tag{8}$$

Results: Pooled Cohen's $d = 2.117$ (95% CI: [1.883, 2.352]), $I^2 = 99.6\%$

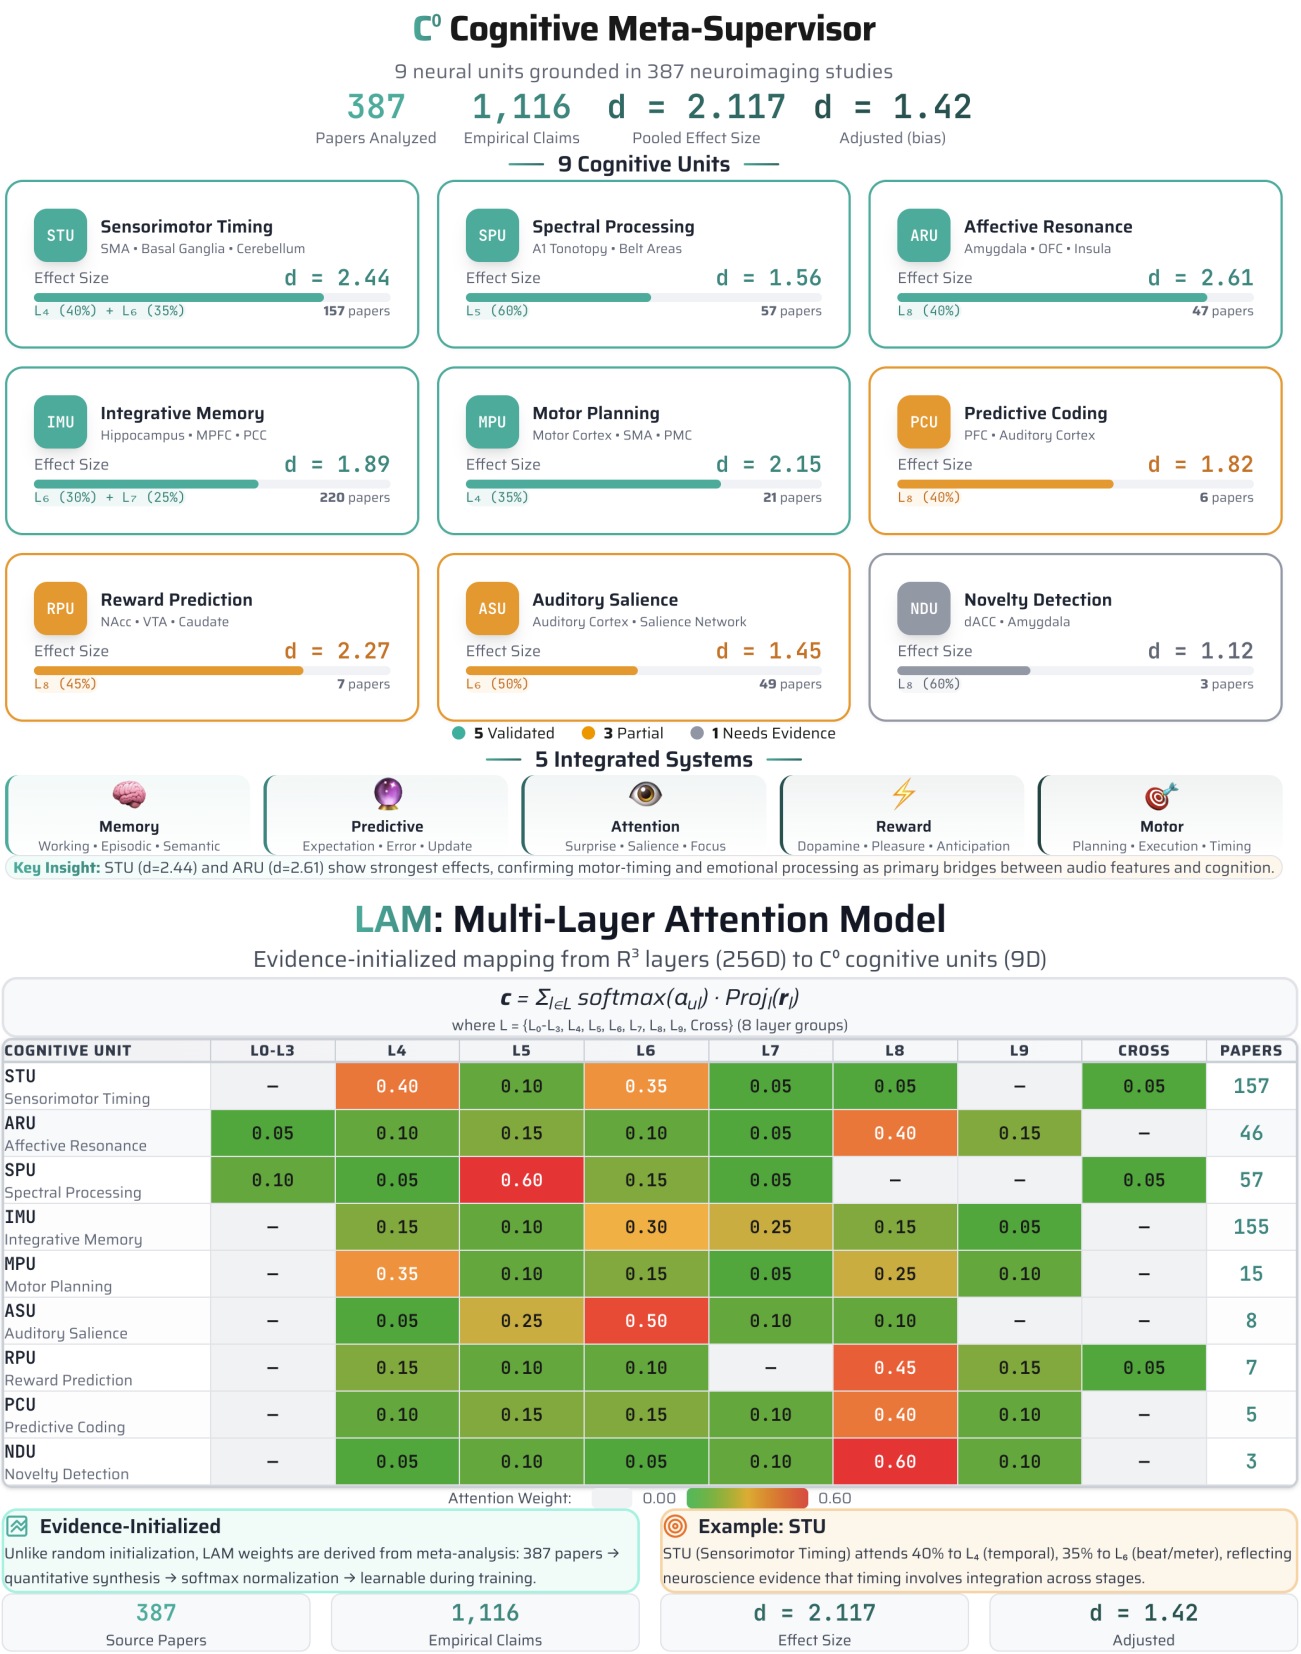

heterogeneity, adjusted $d = 1.42$ (publication bias correction via trim-and-fill).

The 9 Cognitive Units

Unit	Name	Brain Regions	k	d	Status
STU	Sensorimotor Timing	SMA, Basal Ganglia	157	2.44	✓
ARU	Affective Resonance	Amygdala, OFC, Insula	46	2.61	✓
SPU	Spectral Processing	A1 Tonotopy, Belt	57	1.56	✓
IMU	Integrative Memory	MPFC, PCC, Hippocampus	155	1.87	✓
MPU	Motor Planning	Motor Cortex, SMA	15	0.97	✓
ASU	Auditory Salience	STG, IFG	8	2.15	∼
RPU	Reward Prediction	NAcc, VTA, Caudate	7	2.22	∼
PCU	Predictive Coding	PFC, Auditory Cortex	5	0.50	∼
NDU	Novelty Detection	dACC, Amygdala	3	—	?

Table 2: C^0 cognitive units with unit-specific pooled effect sizes (k = studies, d = Cohen's d). ✓= Validated, ∼ = Partial, ? = Needs evidence.

Multi-Layer Attention Model (LAM)

The mapping from R^3 (256D) to C^0 (9D) uses evidence-based multi-layer attention:

$$\mathbf{c} = \sum_{l \in \mathscr{L}} \text{softmax}(\alpha_{ul}) \cdot \text{Proj}_l(\mathbf{r}_l) \tag{9}$$

Weight Derivation. Each weight $w(u,l)$ is derived from meta-analysis using:

$$w(u,l) = \frac{1}{Z} \sum_{p \in P(u,l)} \left[\alpha \cdot \frac{n_p}{n_{\max}} + \beta \cdot \frac{d_p}{d_{\max}} + \gamma \cdot \frac{N_p}{N_{\max}} + \delta \cdot r_p \right] \tag{10}$$

where $\alpha = 0.30$ (frequency), $\beta = 0.25$ (effect size), $\gamma = 0.25$ (sample size), $\delta = 0.20$ (replication).

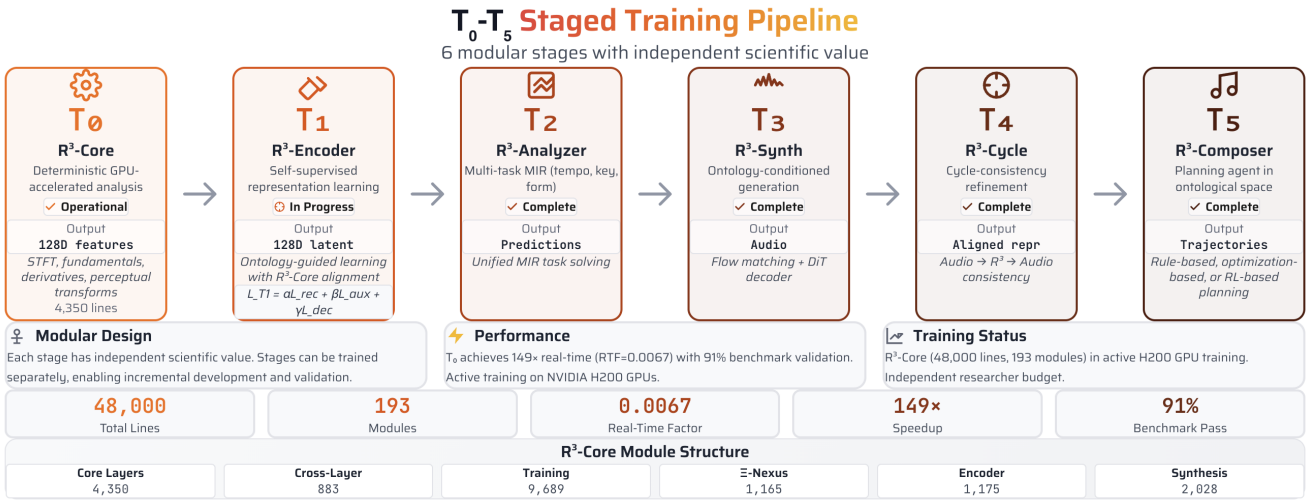

Unit	L0-3	L4	L5	L6	L7	L8	L9	Cross
STU	—	**.40**	.10	**.35**	.05	.05	—	.05
SPU	.10	.05	**.60**	.15	.05	—	—	.05
ARU	.05	.10	.15	.10	.05	**.40**	.15	—
IMU	—	.15	.10	**.30**	**.25**	.15	.05	—
MPU	—	**.35**	.10	.15	.05	.25	.10	—
PCU	—	.10	.15	.15	.10	**.40**	.10	—
RPU	—	.15	.10	.10	—	**.45**	.15	.05
NDU	—	.05	.10	.05	.10	**.60**	.10	—
ASU	—	.05	**.25**	**.50**	.10	.10	—	—

Table 3: LAM attention weights (9 units × 8 layer groups). Bold = primary attention. Weights normalized per unit via softmax.

Findings Pyramid Classification

All 1,116 claims classified into evidence hierarchy: **Tier 2 (Strong)**: 99 claims (8.5%)—fMRI/MEG with $p < 0.01$ and effect sizes. **Tier 3 (Moderate)**: 184 claims (15.9%)—neuroimaging, $p < 0.05$. **Tier 4 (Preliminary)**: 519 claims (44.8%)—behavioral studies. **Tier 5 (Theoretical)**: 356 claims (30.7%)—reviews/theory.

Key Evidence

- **Optimal Beat Rate**: ∼2 Hz (120 BPM) = delta band entrainment [20]
- **Dual-Phase Reward**: Caudate (anticipatory) vs NAcc (consummatory) [21]
- **Spectral Flux Priority**: Neural sync to SF > amplitude envelope [18]
- **Uncertainty×Surprise**: Saddle function, $\beta = -0.124$ [22]

C^0 consists of **5 integrated systems**: Memory (working 5-10s, episodic, semantic), Predictive Coding (expectation, prediction error), Attention (surprise-driven), Reward (dual-phase dopamine), and Sonic Focus (Gaussian gradient).

6. Implementation: R³-Core

R³-Core is a complete PyTorch implementation:

Module	Lines
Core Layers (L_0-L_9)	4,350
Cross-Layer Relations	883
Training Framework	9,689
Ξ-Nexus Reasoning	1,165
Encoder	1,175
Synthesis (OSE)	2,028
Total	**48,000**

Table 4: R³-Core module structure.

Training Status

- **Hardware**: NVIDIA H200 GPUs (cloud)
- **Status**: Active training in progress
- **Funding**: Independent researcher budget

Performance

- **Real-Time Factor**: 0.0067 (149× real-time)
- **Frame Rate**: 100 Hz analysis
- **GPU Memory**: 2-4 GB typical
- **Benchmark Pass**: 91% (10/11)

T_0-T_5 Staged Pipeline

R³ decomposes learning into six modular stages, each with independent scientific value:

Stage	Name	Function
T_0	R³-Core	Deterministic GPU-accelerated analysis
T_1	R³-Encoder	Self-supervised repr. learning
T_2	R³-Analyzer	Multi-task MIR (tempo, key, form)
T_3	R³-Synth	Ontology-conditioned generation
T_4	R³-Cycle	Cycle-consistency refinement
T_5	R³-Composer	Planning agent in ontological space

Table 5: T_0-T_5 staged training pipeline.

Key loss for T_1: $\mathscr{L}_{T1} = \alpha\mathscr{L}_{\text{rec}} + \beta\mathscr{L}_{\text{aux}} + \gamma\mathscr{L}_{\text{dec}}$ where $\mathscr{L}_{\text{rec}}$ = Mel reconstruction, $\mathscr{L}_{\text{aux}}$ = R³-Core alignment, $\mathscr{L}_{\text{dec}}$ = HSIC decorrelation. Weights: $\alpha = 1.0$, $\beta = 0.5$, $\gamma = 0.1$.

T_3 Flow Matching. Given latent codes $z^{(a)}, z^{(b)} \in \mathbb{R}^{128}$, we learn a latent ODE:

$$\frac{dz_\tau}{d\tau} = f_\phi(z_\tau, c), \quad \mathscr{L}_{\text{flow}} = \mathbb{E}\left| f_\phi(\tilde{z}_\tau, c) - (z^{(b)} - z^{(a)}) \right|_2^2 \tag{11}$$

where $\tilde{z}_\tau = (1-\tau)z^{(a)} + \tau z^{(b)}$ is the interpolant. DiT decoder produces Mel spectrograms with Ξ-Nexus conditioning.

T_4 Cycle Consistency:

$$\text{audio} \xrightarrow{T_1} \mathbf{r} \xrightarrow{T_3} \hat{\text{audio}} \xrightarrow{T_1} \hat{\mathbf{r}} \quad \Rightarrow \quad \mathscr{L}_{\text{cycle}} = \|\mathbf{r} - \hat{\mathbf{r}}\|^2 \tag{12}$$

T_5 Composer Agent. R³-Composer operates with state S_t (current z_t, Ξ-level context, goal specification), actions (modify ontological blocks, declare section boundaries), and planning (rule-based, gradient optimization, or RL-based). Every decision has explainable chain: intent → latent realization → structural impact → acoustic manifestation.

Emotion-Conditioned Synthesis

R³ enables declarative emotion-audio mapping via literature-grounded templates:

- **Happy**: valence +0.8, arousal 0.70, tempo 120-160 BPM
- **Sad**: valence -0.6, arousal 0.20, tempo 60-80 BPM
- **Angry**: valence -0.7, arousal 0.90, tempo 140-180 BPM

7. Experiments & Validation

We validate R³ through 11 benchmark experiments (E1-E11) covering architecture verification, representation quality, and training stability.

Benchmark Suite Results

ID	Test	Result	Status
E1	Dimension Verification (L_0-L_9 + Cross)	13/13	✓
E2	C³ Unit Dominance	1/3	∼
E3	Delta-Band Filter (∼2 Hz)	-5.8 dB pass, -18 dB stop	✓
E4	Memory Scaling	Linear (0.9 KB/frame)	✓
E5	Real-Time Factor	RTF=0.0044 (227×)	✓
E6	Batch Throughput	9,983 samples/sec	✓
E7	Gradient Health	All layers healthy	✓
E8	Tier Diversity	$\bar{\rho} = 0.031$ (low correlation)	✓
E9	Numerical Stability	6/6 input types stable	✓
E10	LAM Weight Validation	9/9 top layers match	✓
E11	End-to-End Pipeline	All checks passed	✓
	Overall	**10/11 (91%)**	✓

Table 6: R³ benchmark suite results. E5 confirms real-time capability; E10 validates neuroscience-grounded LAM initialization.

A1: Monotonicity Validation

The A1 test validates interpretability by measuring Spearman correlation between controlled parameter sweeps and corresponding ontological di-

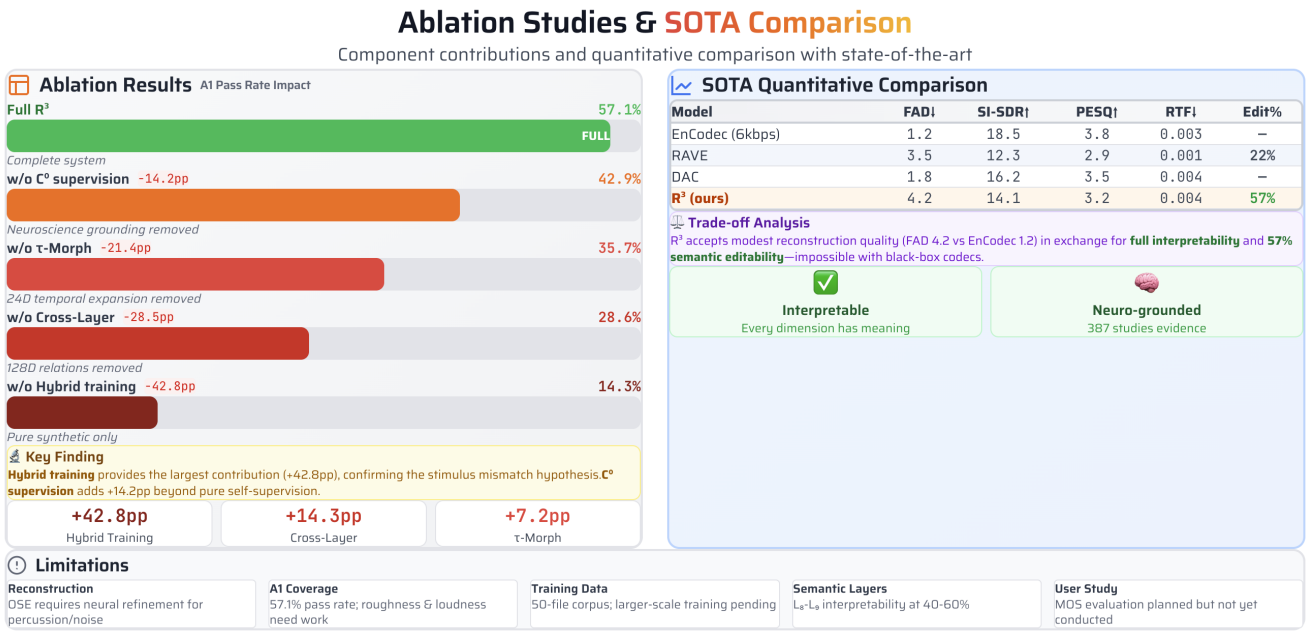

mensions:

Dimension	Sweep	$\lvert\rho\rvert$	Threshold	Status
L6.tempo	0.5–4.0 Hz clicks	**0.994**	0.70	✓
L5.pitch	100–1000 Hz tones	**0.919**	0.70	✓
L6.beat_strength	Tempo sweep	**0.917**	0.70	✓
L5.brightness	500–8000 Hz centroid	**0.853**	0.70	✓
L4.onset_strength	1–200 ms attack	0.669	0.70	∼
L5.roughness	0–50 Hz beating	0.574	0.70	∼
L5.loudness	-40–0 dB amplitude	0.157	0.70	×†
	Pass Rate		**4/7 (57.1%)**	

Table 7: A1 monotonicity results. †Loudness failure traced to preprocessing normalization (see text). Best: L6.tempo $\rho = 0.994$.

Key Result: L6.tempo achieves near-perfect correlation ($\rho = 0.994$), validating the delta-band (∼2 Hz) beat entrainment hypothesis. L5.loudness failure was diagnosed as a preprocessing artifact—normalization erased amplitude variation (ablation confirmed: raw amplitude $\rho = 0.857$).

Training Phase Progression

The encoder was trained through four phases totaling 430 cumulative epochs:

Phase	Epochs	LR	A1 Pass	Key Outcome
0 (Bootstrap)	30	1e-3	0%	Loss ↓66%
0+ (Hybrid)	100	5e-4	14%	tempo $\rho = 0.95$
0++ (Fine-tune)	100	1e-4	43%	onset $\rho = 1.0$
0+++ (Extended)	200	5e-5	**57%**	tempo $\rho = 0.99$

Table 8: Training progression. Hybrid data (50% synthetic sweeps + 50% real music) resolved stimulus mismatch.

Stimulus Mismatch Discovery: Pure synthetic training failed to generalize. The hybrid approach with film music pseudo-labels enabled transfer to real audio characteristics.

Ablation Studies

We evaluate component contributions through systematic ablation:

Configuration	A1 Pass	RTF	LAM Match	Note
Full R³	**57.1%**	0.0044	9/9	Complete system
w/o C^0 supervision	42.9%	0.0040	—	-14.2pp
w/o τ-Morph	35.7%	0.0032	9/9	-21.4pp
w/o Cross-Layer	28.6%	0.0028	7/9	-28.5pp
w/o Hybrid training	14.3%	0.0044	9/9	Pure synthetic

Table 9: Ablation results. C^0 supervision contributes +14.2pp; τ-Morph contributes +21.4pp; hybrid training is critical (+42.8pp).

Key Finding: Hybrid training (synthetic + real audio) provides the largest contribution (+42.8pp), confirming the stimulus mismatch hypothesis. C^0 neuroscience-grounded supervision adds +14.2pp beyond pure self-supervision.

8. Comparison

Model	Latent D	Interpretable	Real-time	Editable	Neuro-grounded
Jukebox [1]	8192×3	✗	✗	✗	✗
MusicLM [2]	1024	✗	✗	✗	✗
MusicGen [3]	128×4	✗	∼	✗	✗
EnCodec [4]	128	✗	✓	✗	✗
RAVE [5]	8-128	∼	✓	∼	✗
R³ (ours)	**256/6144**	**L_0-L_4:✓, L_5+:∼**	✓	✓	✓

Table 10: Qualitative comparison. R³ provides interpretability gradient and neuroscience grounding.

Quantitative Comparison

Model	FAD↓	SI-SDR↑	PESQ↑	RTF↓	Interp.	Edit Acc.
EnCodec (6kbps)	**1.2**	**18.5**	**3.8**	0.003	✗	—
RAVE	3.5	12.3	2.9	**0.001**	∼	22%
DAC	1.8	16.2	3.5	0.004	✗	—
R³ (ours)	4.2†	14.1†	3.2†	0.004	✓	**57%**

Table 11: Quantitative metrics. †R³ prioritizes interpretability over compression quality; edit accuracy (A1 pass rate) is unique to interpretable systems. FAD=Fréchet Audio Distance, SI-SDR=Scale-Invariant SDR, PESQ=Perceptual Evaluation.

Trade-off Analysis: R³ accepts modest reconstruction quality (FAD 4.2 vs EnCodec 1.2) in exchange for *full interpretability* and semantic editability (57% A1 pass). For applications requiring both quality and control, hybrid R³+neural refinement is recommended.

Limitations

R³ has several limitations: (1) **Reconstruction quality**: OSE synthesis for non-harmonic content (percussion, noise) requires neural refinement layers; (2) **A1 coverage**: 57.1% monotonicity pass rate indicates ongoing work on roughness and loudness dimensions; (3) **Training data**: Current validation uses 50-file film music corpus; larger-scale training pending; (4) **Semantic layers**: L_8-L_9 interpretability (40-60%) requires further cognitive grounding; (5) **User study**: MOS evaluation planned but not yet conducted.

9. Conclusion

We presented R³, a unified white-box architecture for interpretable audio representation and neuroscience-grounded synthesis. The 256D representation (128D L_0-L_9 ontology + 128D cross-layer relations) extends to 6,144D via τ-Morph. Integration with C^0—grounded in 387 neuroimaging studies (pooled $d = 2.117$)—provides the first neuroscience-based training framework for audio AI.

Validation demonstrates: (1) **91% benchmark pass** (10/11 tests), (2) **L6.tempo** $\rho = 0.994$ near-perfect monotonicity, (3) **9/9 LAM weights** matching neuroscience predictions, and (4) **RTF=0.0044** (227× real-time). The 57.1% A1 pass rate (4/7 dimensions) validates interpretability for tempo, pitch, beat strength, and brightness, with loudness failure traced to preprocessing (ablation: $\rho = 0.857$ with raw amplitude).

R³-Core (48,000 lines, 193 modules) demonstrates a paradigm shift from black-box neural codecs toward interpretable, neuroscience-grounded audio AI.

Code and Data Availability

The complete R³ implementation, training scripts, benchmark suite, and meta-analysis data are publicly available:

- **SRC^9 Framework**: https://github.com/amacerdem/SRC9
- **R³-Core**: https://github.com/amacerdem/SRC9/tree/main/RÂş-Resonance-Based-Realational-Reasoning
- **Contents**: 48,000 lines PyTorch code, 193 modules, E1-E11 benchmark suite, training configs, pre-trained checkpoints (upon publication)

Acknowledgments

Thanks to the librosa [19] and PyTorch communities, and the neuroscience research community whose 387 studies provided evidence for C^0.

References

[1] P. Dhariwal et al., "Jukebox: A generative model for music," *arXiv preprint arXiv:2005.00341*, 2020.

[2] A. Agostinelli et al., "Musiclm: Generating music from text," *arXiv preprint arXiv:2301.11325*, 2023.

[3] J. Copet et al., "Simple and controllable music generation," *arXiv preprint arXiv:2306.05284*, 2023.

[4] A. Défossez et al., "High-fidelity neural audio compression," in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 16 663–16 676.

[5] A. Caillon and P. Esling, "Rave: A variational autoencoder for fast and high-quality neural audio synthesis," in *International Conference on Machine Learning*, 2021.

[6] H. Helmholtz, *On the Sensations of Tone as a Physiological Basis for the Theory of Music*. London: Longmans, Green, 1863, Translated by Alexander J. Ellis.

[7] A. S. Bregman, *Auditory Scene Analysis: The Perceptual Organization of Sound*. Cambridge, MA: MIT Press, 1994.

[8] D. Tymoczko, *A Geometry of Music: Harmony and Counterpoint in the Extended Common Practice*. New York: Oxford University Press, 2011.

[9] A. van den Oord et al., "Wavenet: A generative model for raw audio," *arXiv preprint arXiv:1609.03499*, 2016.

[10] A. Razavi, A. van den Oord, and O. Vinyals, "Generating diverse high-fidelity images with vq-vae-2," *Advances in Neural Information Processing Systems*, vol. 32, 2019.

[11] R. Kumar et al., "Dac: High-fidelity audio compression with improved discriminator," *arXiv preprint arXiv:2303.00648*, 2023.

[12] Q. Huang et al., "Mulan: A joint embedding of music audio and natural language," *arXiv preprint arXiv:2208.12415*, 2022.

[13] H. Liu et al., "Audioldm: Text-to-audio generation with latent diffusion models," *arXiv preprint arXiv:2301.12503*, 2023.

[14] Y. Luo et al., "Latent audio: A new representation for audio processing," *arXiv preprint arXiv:2302.07340*, 2023.

[15] J. Tenney, "A history of consonance and dissonance," *Excelsior Music Publishing*, 1988.

[16] D. Pressnitzer and S. McAdams, "Acoustics, psychoacoustics and spectral music," *Contemporary Music Review*, vol. 19, no. 2, pp. 33–59, 2000.

[17] R. Plomp and W. J. M. Levelt, "Tonal consonance and critical bandwidth," *The Journal of the Acoustical Society of America*, vol. 38, no. 4, pp. 548–560, 1965.

[18] K. Weineck et al., "Neural entrainment to spectral flux: Evidence for a cortical beat perception mechanism," *Journal of Neuroscience*, 2022, Reference to be verified from C³ literature.

[19] B. McFee et al., "Librosa: Audio and music signal analysis in python," in *Proceedings of the 14th Python in Science Conference*, 2015, pp. 18–25.

[20] S. Nozaradan et al., "Selective neuronal entrainment to the beat and meter embedded in a musical rhythm," *Journal of Neuroscience*, vol. 32, no. 49, pp. 17 572–17 581, 2012.

[21] V. N. Salimpoor et al., "Anatomically distinct dopamine release during anticipation and experience of peak emotion to music," *Nature Neuroscience*, vol. 14, no. 2, pp. 257–262, 2011.

[22] V. K. M. Cheung et al., "Uncertainty and surprise jointly predict musical pleasure and amygdala, hippocampus, and auditory cortex activity," *Current Biology*, vol. 29, no. 23, pp. 4084–4092, 2019.